

Assignment: Build a Fact-Check AI Agent

Objective

Build an AI **Fact-Check Agent** using **LangChain**, **FastAPI**, **FastMCP**, and **Gemini API**.
The agent must verify factual claims by **calling MCP tools** and returning a reasoned verdict.

Mandatory MCP Tools (Both Required)

MCP Tool 1: Web Search Tool

- Implemented using **FastMCP**
 - Uses an **external web search API** (e.g. SerpAPI)
 - Purpose:
 - Retrieve general web evidence related to a claim
-

MCP Tool 2: News Search Tool

- Implemented using **FastMCP**
 - Uses a **news API** (e.g. NewsAPI)
 - Purpose:
 - Retrieve **recent or authoritative news articles**
-

Agent Behavior

1. Accept a **factual claim** from the user

2. Use LangChain reasoning to:

- Decide search queries
- Call **both MCP tools**

3. Combine results from:

- Web search
- News search

4. Analyze evidence using Gemini

5. Return a structured verdict

Verdict Classification (Required)

The agent must classify each claim as one of:

-  **Likely True**
-  **Uncertain**
-  **Likely False**

Each verdict **must include justification** based on tool results.

API Specification

Endpoint

POST /fact-check

Input

{

```
        "claim": "string"
    }
```

Output

```
{
    "verdict": "Likely True | Uncertain | Likely False",
    "explanation": "string",
    "sources": [
        {
            "title": "string",
            "url": "string",
            "type": "web | news"
        }
    ],
    "tools_used": [ "web_search", "news_search" ]
}
```

Technical Requirements

- Python
- LangChain ReAct agent (tool-calling required)
- FastAPI for API exposure
- FastMCP for **all MCP tools**
- Gemini API as LLM backend
- `.env` file for API keys
- No hardcoded secrets

Test Claims (They **MUST** Work)

- “The Earth is flat.”
- “OpenAI released GPT-4 in 2023.”
- “Drinking bleach can cure COVID-19.”
- “Google is shutting down Gmail in 2025.”

If the agent answers these **without MCP calls**, it fails.

Evaluation Criteria

- Proper MCP tool implementation
 - Evidence-based reasoning (not hallucination)
 - Clean LangChain agent logic
 - Clear verdict explanation
 - Code clarity & structure
-

Hard Rule

“If your agent produces a verdict without calling **both MCP tools**, the assignment is automatically rejected.”